

Scam Honeypot — ForgeHacks 2026 Demo Script (2–3 min)

Track 05: AI + Cybersecurity

Total runtime target: ~3:00. Timestamps are approximate — practice once with the app warmed up.

Before you record

1. **Warm up the model** — run one conversation first. The first LLM call after idle has a ~20s cold start.
 2. **Have the deployed link ready** — your live Streamlit Cloud app URL (used in the 2:20 beat).
 3. **Never show your API key** — no terminals, no `.env`, no Featherless dashboard on screen.
 4. Record the browser tab only, 1920×1080, with your mic.
 5. The live run takes ~5–10 min; you'll **narrate over it and edit out the dead air** (or pre-run, then screen-record the finished conversation with the "Replay cached demo" mode).
 6. Plan B if the API flakes: **Offline simulator** mode — instant, no API calls.
-

0:00–0:15 — Intro

On screen: Title card: "Scam Honeypot — ForgeHacks 2026, Track 05: AI + Cybersecurity" + your name/team.

Say:

"Hi, I'm [name], and this is Scam Honeypot — an AI-powered decoy that wastes scammers' time and harvests their fraud infrastructure for law enforcement. Here's the problem it solves."

0:15–0:35 — Problem

On screen: Keep title card or show the app's landing page (red SIMULATED banner visible).

Say:

"Every day in India, thousands of people — most of them elderly — lose their life savings to phone scams: fake police arrests, KYC fraud, courier threats. In 2024 alone, Indians lost over ₹11,000 crore to cyber fraud. And it's getting worse: scammers now use AI to make thousands of calls a day. They win because they're fast, automated, and relentless. So we asked: what if AI fought back — at the same scale?"

0:35–1:10 — Setup & concept

On screen: Sidebar. Point at each control as you mention it.

Say:

"Meet Scam HoneyPot — a defensive AI tool that wastes scammers' time and harvests their payment details for law enforcement. Here's the idea: instead of hanging up, we put a decoy on the line — an AI persona, Ramesh Gupta, a 62-year-old retired government employee. Ramesh sounds confused and scared, so the scammer stays on the line... while the system quietly extracts every fraud indicator the scammer reveals."

Do: - Pick **Digital Arrest** from the Scam Script dropdown. - Point out Mode is **Live (LLM)** — "two real LLM agents talking to each other, powered by Qwen2.5-72B through Featherless." - Point at the red banner: "**This is a simulated demo. No real scammer, no real data — safety first.**" - Click ► **Start**.

1:10–1:55 — The live conversation

On screen: Conversation bubbles appear (👤 scammer / 😊 persona), Time Wasted ticker running.

Say (as the first messages land):

"The scammer agent opens with a classic digital-arrest script — 'courier intercepted in your name, five hundred grams of illegal substances.' Ramesh doesn't accuse or hang up — he stalls. Every reply ends with a confused question that makes the scammer repeat themselves, because repetition is exactly what we want: the more the scammer talks, the more evidence they leave behind."

Say (mid-run, as more turns land):

"And here's the critical part: Ramesh can only ever reveal fake 'canary' data — nothing real. A guardrail layer checks every single reply before it's sent. If the model ever tries to leak a real value, the reply is blocked and rewritten. The Leaks Blocked counter stays at zero — by design."

Say (as scammer escalates — link, UPI ID, account number):

"Watch what happens when the scammer gets impatient — he escalates to the actual money step: a fake RBI verification link... a UPI ID... a bank account number. Every one of these is being captured live by the extractor."

1:55–2:20 — The payoff: intel & report

On screen: Scroll to **Collected Intel** panel, then **Generated Report**.

Say:

"When the call finally ends, the payoff appears: the scam type, the impersonated identity, the UPI ID, the phone number, the phishing URL, the bank account — every indicator the scammer burned time revealing. And it's not just a list — the system generates a full report: a summary of the interaction, the scam stages it observed — impersonation, fear hook, isolation, urgency, extraction — and a draft complaint, pre-formatted for the Cybercrime Helpline 1930."

Do: - Point at **Scam Script Stages Observed**. - Click **↓ Download report.md** — "one click, and the victim has an evidence packet ready to file."

2:20–2:40 — Why it matters + live deployment

On screen: Browser switches to (or shows the URL of) your deployed Streamlit app.

Say:

"This isn't just a cool demo. Every minute a scammer spends on Ramesh is a minute they're not scamming a real victim. The indicators we collect — UPI IDs, phone numbers, phishing domains — can be fed straight into bank blocklists and caller-protection systems. And the report gives cybercrime cells exactly what they need to act on a complaint. Best part: it's already live. The whole thing is deployed on Streamlit Community Cloud — free to host, and the app reads its secrets from the platform, so my API key never touches the codebase. Anyone with the link can run a simulation right now."

2:40–3:00 — Architecture + outro

On screen: Keep the report visible; optionally a simple architecture diagram.

Say:

"Under the hood: two LLM agents — a scammer simulator and a decoy persona — a guardrail that enforces the canary rules, an extractor that pulls indicators out of raw chat, and a Streamlit dashboard on top. Everything is configurable — four scam scripts, live or offline modes. Scam HoneyPot flips the economics of fraud: the scammer's biggest advantage is time, and now their time is being stolen by a machine that never gets tired, never gets scared, and files the evidence for free. Thank you."

Editing checklist

- Cut all LLM wait time between turns; keep the flow tight.
- Keep the red SIMULATED banner visible the whole time — it's the ethics story.
- End frame: the report + "Download report.md" button (or your team name/title card).
- Export 1080p, H.264, ≤3 min per ForgeHacks spec.